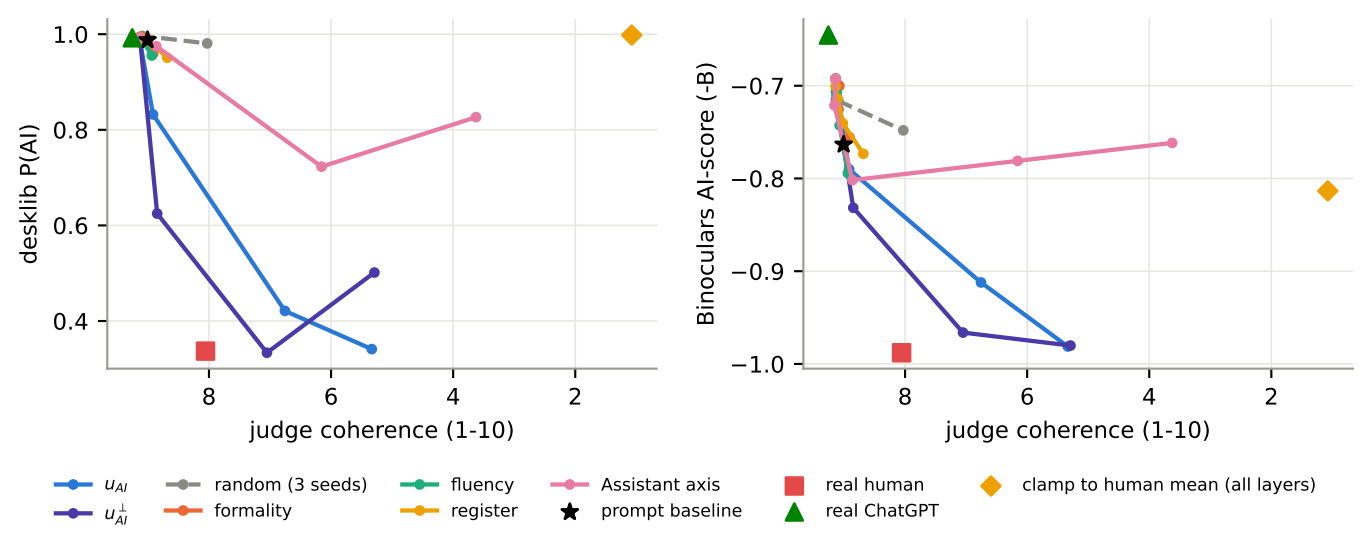
desklib P(AI)
1.0
0.8
0.6
0.4
8
6
4
2
judge coherence (1-10)
Binoculars AI-score (-B)
−0.7
−0.8
−0.9
−1.0
8
6
4
2
judge coherence (1-10)
u_{AI}
$u_{AI}^{\perp}$
random (3 seeds)
formality
fluency
register
Assistant axis
prompt baseline
real human
real ChatGPT
clamp to human mean (all layers)